

Cloud Computing

Containers and Virtual Machines

Charlotte Laclau

Télécom Saint-Etienne



Outline of today's course

- ① Quizz : do you remember something ?
- ② Course Part 1 : Containers vs. Virtual Machines
- ③ Course Part 2 : Data security, Infrastructure Management and Service Level Agreement
- ④ Lab 1 follow-up



Do you remember something ?

What is Cloud Computing?

- ① A software package
- ② A spreadsheet
- ③ A way to store data remotely
- ④ A technique for managing computer applications in companies
- ⑤ An OLAP cube



Do you remember something ?

What is Cloud Computing?

- ① A software package
- ② A spreadsheet
- ③ A way to store data remotely
- ④ A technique for managing computer applications in companies
- ⑤ An OLAP cube



Do you remember something ?

What are the real advantages of Cloud Computing?

- ① Strong security
- ② Outsourcing of the IT department
- ③ Cost reduction
- ④ Easy access
- ⑤ Infinite capacity



Do you remember something ?

What are the real advantages of Cloud Computing?

- ① Strong security
- ② Outsourcing of the IT department
- ③ Cost reduction
- ④ Easy access
- ⑤ Infinite capacity



Do you remember something ?

What are the disadvantages of Cloud Computing?

- ① A limited offer
- ② Security breaches
- ③ Strong security
- ④ Dependence on its supplier
- ⑤ The price of subscriptions
- ⑥ The cost of managing the cloud's IT infrastructure



Do you remember something ?

What are the disadvantages of Cloud Computing?

- ① A limited offer
- ② Security breaches
- ③ Strong security
- ④ Dependence on its supplier
- ⑤ The price of subscriptions
- ⑥ The cost of managing the cloud's IT infrastructure



Do you remember something ?

What is a public cloud?

- ① A cloud formation that can be seen across the globe
- ② A cloud service that can only be accessed from a publicly shared computer
- ③ A multi-tenant cloud environment accessed over the internet
- ④ A cloud environment owned, operated and controlled by a public company



Do you remember something ?

What is a public cloud?

- ① A cloud formation that can be seen across the globe
- ② A cloud service that can only be accessed from a publicly shared computer
- ③ A multi-tenant cloud environment accessed over the internet
- ④ A cloud environment owned, operated and controlled by a public company



Do you remember something ?

Which of the following statements could be used to describe a private cloud deployment?

- ① A cloud environment maintained within an enterprise's own data center
- ② A private environment within a third-party or public cloud provider's architecture
- ③ A cloud service inaccessible to anyone but the cultural elite
- ④ Both #1 and #2



Do you remember something ?

Which of the following statements could be used to describe a private cloud deployment?

- ① A cloud environment maintained within an enterprise's own data center
- ② A private environment within a third-party or public cloud provider's architecture
- ③ A cloud service inaccessible to anyone but the cultural elite
- ④ Both #1 and #2



Do you remember something ?

True or false: Open source tools can be used to manage private and public clouds.

- ① FALSE
- ② TRUE



Do you remember something ?

True or false: Open source tools can be used to manage private and public clouds.

- ① FALSE
- ② TRUE



Do you remember something ?

Amazon Elastic Compute Cloud (EC2) is primarily considered which type of cloud computing model?

- ① Software as a service (SaaS)
- ② Platform as a service (PaaS)
- ③ Infrastructure as a service (IaaS)
- ④ None of the above



Do you remember something ?

Amazon Elastic Compute Cloud (EC2) is primarily considered which type of cloud computing model?

- ① Software as a service (SaaS)
- ② Platform as a service (PaaS)
- ③ Infrastructure as a service (IaaS)
- ④ None of the above



Do you remember something ?

True or false: A multi-cloud strategy is the same as hybrid cloud.

- ① TRUE
- ② FALSE



Do you remember something ?

True or false: A multi-cloud strategy is the same as hybrid cloud.

- ① TRUE
- ② FALSE



Do you remember something ?

True or false: Enterprises can physically ship their data to some public cloud providers to upload that data into the cloud.

- ① FALSE
- ② TRUE



Do you remember something ?

True or false: Enterprises can physically ship their data to some public cloud providers to upload that data into the cloud.

① FALSE

② TRUE



Do you remember something ?

The Cloud Computing will end the existence of the direction of the computer systems in company if :

- ① The Cloud computing corporation go public
- ② The costs and the quality of the services are competitive
- ③ Confidentiality of the data is assured




Do you remember something ?

The Cloud Computing will end the existence of the direction of the computer systems in company if :

- ① The Cloud computing corporation go public
- ② The costs and the quality of the service are competitive
- ③ Confidentiality of the data is assured




Overview



Virtualization

Definition: running on the same physical machine, multiple systems as if they were running on separate physical machines

- ▶ a technology that transforms hardware into software
- ▶ allows to run multiple OS as **virtual machines**



Virtual Machines (VM's)

- ▶ **Virtual Machine** : an emulation of a real computer that executes programs like a real computer.



Virtual Machines (VM's)

- ▶ **Virtual Machine** : an emulation of a real computer that executes programs like a real computer.
- ▶ VM's run on top of a physical machine using a **hypervisor**



Virtual Machines (VM's)

- ▶ **Virtual Machine** : an emulation of a real computer that executes programs like a real computer.
- ▶ VM's run on top of a physical machine using a **hypervisor**
- ▶ **Hypervisor** : a tool which control and distributes the hardware resources to each OS according to their needs




Virtual Machines (VM's)

- ▶ **Virtual Machine** : an emulation of a real computer that executes programs like a real computer.
- ▶ VM's run on top of a physical machine using a **hypervisor**
- ▶ **Hypervisor** : a tool which control and distributes the hardware resources to each OS according to their needs

Roles of a hypervisor

- ▶ provides control of the processor and the resources of the host machine
- ▶ allocate to each virtual machine the resources it needs
- ▶ make sure that VM's do not interfere with each other



Host Metal Hypervisor (type 2)

- ▶ A hosted virtualization hypervisor runs on the operating system of the host machine




Host Metal Hypervisor (type 2)

- ▶ A hosted virtualization hypervisor runs on the operating system of the host machine
- ▶ Example :
 - ▶ A computer running OSX can have a VM (e.g. VirtualBox or VMware Workstation 8) installed on top of that OS.
 - ▶ The VM doesn't have direct access to hardware, so it has to go through the host operating system (in our case, the Mac's OSX).




Host Metal Hypervisor (type 2)

- ▶ A hosted virtualization hypervisor runs on the operating system of the host machine
- ▶ Example :
 - ▶ A computer running OSX can have a VM (e.g. VirtualBox or VMware Workstation 8) installed on top of that OS.
 - ▶ The VM doesn't have direct access to hardware, so it has to go through the host operating system (in our case, the Mac's OSX).



Host Metal Hypervisor (type 2)

- ▶ A hosted virtualization hypervisor runs on the operating system of the host machine
- ▶ Example :
 - ▶ A computer running OSX can have a VM (e.g. VirtualBox or VMware Workstation 8) installed on top of that OS.
 - ▶ The VM doesn't have direct access to hardware, so it has to go through the host operating system (in our case, the Mac's OSX).




+ more hardware compatibility.



Host Metal Hypervisor (type 2)

- ▶ A hosted virtualization hypervisor runs on the operating system of the host machine
- ▶ Example :
 - ▶ A computer running OSX can have a VM (e.g. VirtualBox or VMware Workstation 8) installed on top of that OS.
 - ▶ The VM doesn't have direct access to hardware, so it has to go through the host operating system (in our case, the Mac's OSX).



- + more hardware compatibility.
- lower performance of the VM




Bare Metal Hypervisor (type 1)

- ▶ A bare metal hypervisor is directly executed on the hardware.
 - ▶ Interfaces directly with the underlying hardware
 - ▶ Does not need a host OS to run on
 - ▶ The first thing installed on a host machine's server as the OS is the hypervisor




Bare Metal Hypervisor (type 1)

- ▶ A bare metal hypervisor is directly executed on the hardware.
 - ▶ Interfaces directly with the underlying hardware
 - ▶ Does not need a host OS to run on
 - ▶ The first thing installed on a host machine's server as the OS is the hypervisor



Bare Metal Hypervisor (type 1)

- ▶ A bare metal hypervisor is directly executed on the hardware.
 - ▶ Interfaces directly with the underlying hardware
 - ▶ Does not need a host OS to run on
 - ▶ The first thing installed on a host machine's server as the OS is the hypervisor




+ better performance, scalability, and stability



Bare Metal Hypervisor (type 1)

- ▶ A bare metal hypervisor is directly executed on the hardware.
 - ▶ Interfaces directly with the underlying hardware
 - ▶ Does not need a host OS to run on
 - ▶ The first thing installed on a host machine's server as the OS is the hypervisor




- + better performance, scalability, and stability
- hardware compatibility is limited



Virtual Machines and Hypervisor

Why do we need this additional hypervisor layer in between the VM and the host machine?



Examples of Bare-metal Virtualization Hypervisors (1)

VMware ESX and ESXi

- ▶ Belongs to VMware
- ▶ Mature & Stable tool
- ▶ Free edition available with limited features
- ▶ 5 commercial editions
- ▶ Includes its own kernel



Examples of Bare-metal Virtualization Hypervisors (1)

VMware ESX and ESXi

- ▶ Belongs to Vmware
- ▶ Mature & Stable tool
- ▶ Free edition available with limited features
- ▶ 5 commercial editions
- ▶ Includes its own kernel



Microsoft Hyper-V

- ▶ Belongs to Microsoft
- ▶ Good for Small-Medium Businesses
- ▶ Good for running windows
- ▶ Free edition available with limited features
- ▶ 4 commercial editions



Examples of Bare-metal Virtualization Hypervisors (2)

Citrix XenServer

- ▶ Started as an open source project
- ▶ Mature & Stable tool
- ▶ Free edition available with limited features
- ▶ 2 commercial editions
- ▶ One of the cheapest



Examples of Bare-metal Virtualization Hypervisors (2)

Citrix XenServer

- ▶ Started as an open source project
- ▶ Mature & Stable tool
- ▶ Free edition available with limited features
- ▶ 2 commercial editions
- ▶ One of the cheapest



Oracle VM

- ▶ Belongs to Oracle Corporation
- ▶ Homegrown hypervisor tech. based on Xen
- ▶ Simple, no-frills hypervisor
- ▶ Lacks advanced features



Examples of Hosted Virtualization Hypervisors (1)

VMware Workstation/Fusion/Player

- ▶ Good for running multiple different OS or versions of one OS on one desktop
- ▶ Good for labs and demonstration objectives
- ▶ Good Option for running Windows and Linux on Macs



Examples of Hosted Virtualization Hypervisors (1)

VMware Workstation/Fusion/Player

- ▶ Good for running multiple different OS or versions of one OS on one desktop
- ▶ Good for labs and demonstration objectives
- ▶ Good Option for running Windows and Linux on Macs



Oracle VM VirtualBox

- ▶ Belongs to Oracle Corporation
- ▶ Open source hypervisor
- ▶ Similar to VMware and Microsoft Hyper-V
- ▶ Mature and Stable
- ▶ Suitable for Small-Medium Business & Enterprises



Examples of Hosted Virtualization Hypervisors (2)

KVM

- ▶ **Open Source**
- ▶ Belongs to Red Hat, Inc.
- ▶ Contains features of Type 1 & Type 2
- ▶ Compatible with numerous guest operating systems work with KVM



Examples of Hosted Virtualization Hypervisors (2)

KVM

- ▶ **Open Source**
- ▶ Belongs to Red Hat, Inc.
- ▶ Contains features of Type 1 & Type 2
- ▶ Compatible with numerous guest operating systems work with KVM



Parallels Desktop

- ▶ Belongs to Parallels
- ▶ Good for running Windows on OS
- ▶ 3 commercial editions



How should I choose a hypervisor ?

- ▶ Understand your needs, i.e., flexibility, scalability, reliable support etc.
- ▶ Understand the features, i.e., live migration, storage migration, dynamic memory, automated workflows, etc.
- ▶ Investigate the ecosystem
 - ▶ Possible to evaluate every virtualization hypervisor for free
- ▶ Compare costs
- ▶ **Remember** : VMs are portable and easily convertible



Containers

- ▶ A container provides operating-system-level virtualization





Containers

- ▶ A container provides operating-system-level virtualization
- ▶ Share similarities with a VM
 - ▶ they have private space for processing
 - ▶ can execute commands as root
 - ▶ have a private network interface and IP address
 - ▶ can mount file systems



Containers

- ▶ A container provides operating-system-level virtualization
- ▶ Share similarities with a VM
 - ▶ they have private space for processing
 - ▶ can execute commands as root
 - ▶ have a private network interface and IP address
 - ▶ can mount file systems
- ▶ **Main difference** : containers **share** the host system's kernel with other containers



What is Docker ?

- ▶ An open-source project based on Linux containers.




What is Docker ?

- ▶ An open-source project based on Linux containers.
- ▶ Uses Linux Kernel features to create containers on top of an OS.



What is Docker ?

- ▶ An open-source project based on Linux containers.
- ▶ Uses Linux Kernel features to create containers on top of an OS.
- ▶ A gentle introduction



Why using Docker ?



Why using Docker ?

► Ease of use

- can quickly build and test portable applications
- allows anyone to package an application on their laptop
- *build once, run anywhere*



Why using Docker ?

► **Ease of use**

- can quickly build and test portable applications
- allows anyone to package an application on their laptop
- *build once, run anywhere*

► **Speed** : lightweight and fast

- Take up fewer resources
- No need to boot up a full virtual OS every time



Why using Docker ?

► **Ease of use**

- can quickly build and test portable applications
- allows anyone to package an application on their laptop
- *build once, run anywhere*

► **Speed** : lightweight and fast

- Take up fewer resources
- No need to boot up a full virtual OS every time

► **Docker Hub** : an increasingly rich ecosystem

- an *App store* for Docker images
- public images created by the community



Why using Docker ?

► **Ease of use**

- can quickly build and test portable applications
- allows anyone to package an application on their laptop
- *build once, run anywhere*

► **Speed** : lightweight and fast

- Take up fewer resources
- No need to boot up a full virtual OS every time

► **Docker Hub** : an increasingly rich ecosystem

- an *App store* for Docker images
- public images created by the community

► **Modularity and Scalability** :


- Easy to link containers together to create an app
- Easy to scale or update components independently in the future



Fundamental Concepts of Docker

Docker engine is the layer on which Docker runs.


- ▶ A **Docker Daemon** that runs in the host computer.
- ▶ A **Docker Client** that then communicates with the Docker Daemon to execute commands.
- ▶ Can run on the same machine or on two different ones



Fundamental Concepts of Docker

Docker engine is the layer on which Docker runs.

- ▶ A **Docker Daemon** that runs in the host computer.
- ▶ A **Docker Client** that then communicates with the Docker Daemon to execute commands.
- ▶ Can run on the same machine or on two different ones




Dockerfile : where you write the instructions to build a Docker image.



Fundamental Concepts of Docker

Docker engine is the layer on which Docker runs.

- ▶ A **Docker Daemon** that runs in the host computer.
- ▶ A **Docker Client** that then communicates with the Docker Daemon to execute commands.
- ▶ Can run on the same machine or on two different ones



Dockerfile : where you write the instructions to build a Docker image.

Docker Image : read-only templates built from a set of instructions written in your Dockerfile.



What is Kubernetes?



- ▶ Kubernetes and Docker both help you run applications within linux containers



What is Kubernetes?



- ▶ Kubernetes and Docker both help you run applications within linux containers
- ▶ What if I want to run multiple containers across multiple machines ?
 - ▶ Need to start the right containers at the right time,
 - ▶ Figure out how they can talk to each other
 - ▶ Handle storage considerations
 - ▶ Deal with failed containers or hardware.



What is Kubernetes?



- ▶ Kubernetes and Docker both help you run applications within linux containers
- ▶ What if I want to run multiple containers across multiple machines ?
 - ▶ Need to start the right containers at the right time,
 - ▶ Figure out how they can talk to each other
 - ▶ Handle storage considerations
 - ▶ Deal with failed containers or hardware.
- ▶ **Kubernetes is an open source container orchestration platform**
 - ▶ Running containers across many different machines
 - ▶ Scaling up or down by adding or removing containers when demand changes
 - ▶ Keeping storage consistent with multiple instances of an application
 - ▶ Distributing load between the containers
 - ▶ Launching new containers on different machines if something fails



Benefits of virtualization

- ▶ For a company
 - ▶ Removal of special hardware and utility requirements
 - ▶ Effective management of resources
 - ▶ Better accessibility
 - ▶ Reduced risk of data loss
- ▶ For Data Centers
 - ▶ Maximization of server capabilities
 - ▶ Smaller footprint



Summing-up

Benefits of virtualization

- ▶ For a company
 - ▶ Removal of special hardware and utility requirements
 - ▶ Effective management of resources
 - ▶ Better accessibility
 - ▶ Reduced risk of data loss
- ▶ For Data Centers
 - ▶ Maximization of server capabilities
 - ▶ Smaller footprint

Connection to Cloud Computing

Virtualization in **cloud computing** allows you to run multiple applications and OS's on the same server, thereby providing for **efficient resource utilization** and reducing costs.



Deploying VM's and Containers with AWS

- ▶ VM's ? You already know how to do it !



Deploying VM's and Containers with AWS

- ▶ VM's ? You already know how to do it !
- ▶ Containers ? Can be broken down into 3 categories



Deploying VM's and Containers with AWS

- ▶ VM's ? You already know how to do it !
- ▶ Containers ? Can be broken down into 3 categories
 - ▶ Registry : AWS Elastic Container Registry (ECR)



Deploying VM's and Containers with AWS

- ▶ VM's ? You already know how to do it !
- ▶ Containers ? Can be broken down into 3 categories
 - ▶ Registry : AWS Elastic Container Registry (ECR)
 - ▶ Orchestration: Amazon Elastic Container Service and Amazon Elastic Kubernetes Service

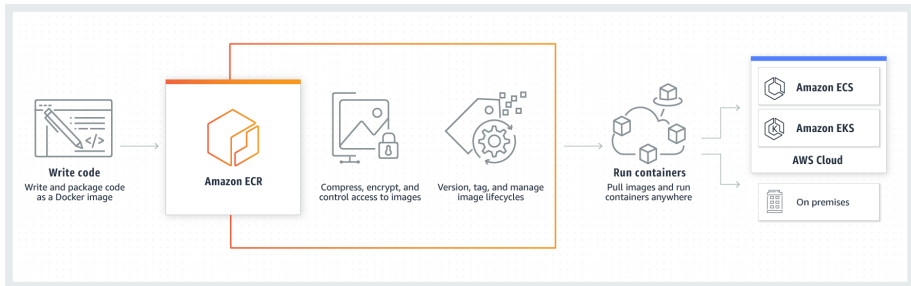


Deploying VM's and Containers with AWS

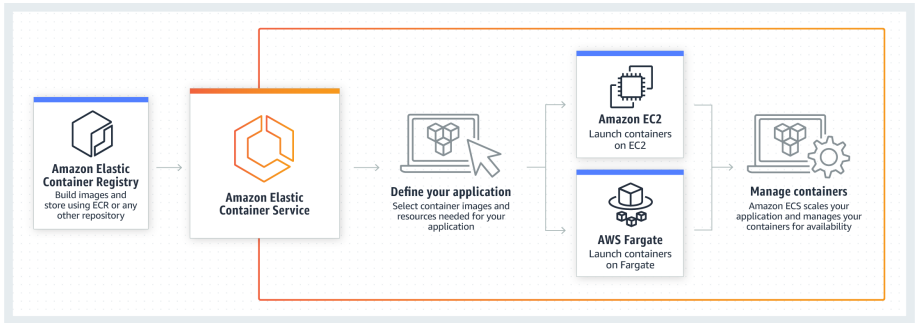
- ▶ VM's ? You already know how to do it !
- ▶ Containers ? Can be broken down into 3 categories
 - ▶ Registry : AWS Elastic Container Registry (ECR)
 - ▶ Orchestration: Amazon Elastic Container Service and Amazon Elastic Kubernetes Service
 - ▶ Compute: AWS Fargate and Amazon EC2



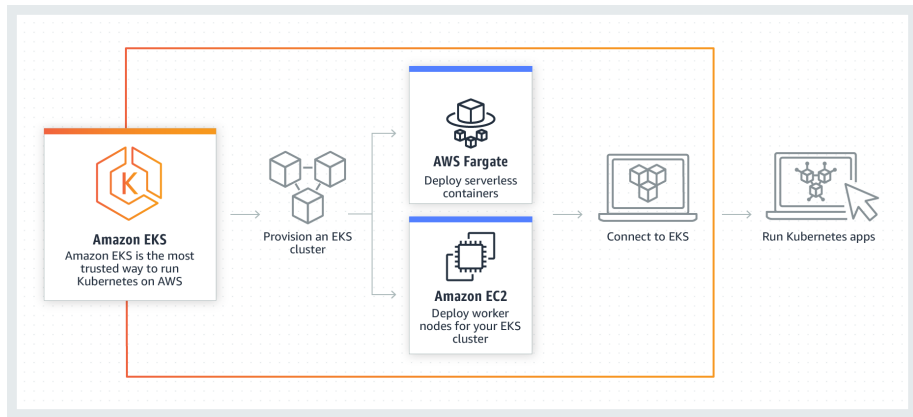
Registry: ECR



Orchestration: ECS

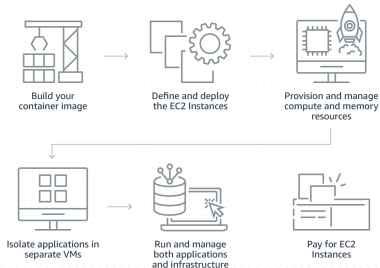


Orchestration: EKS

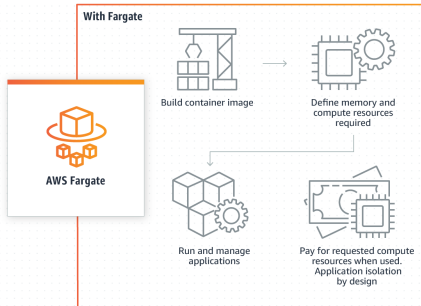


Compute: EC2 / Fargate

Without Fargate



With Fargate



Part 2

Data Security, Infrastructure Management and Service Level Agreement



Data Security in the cloud

- ▶ Similar to protecting data within a traditional data center



- ▶ Responsibility is shared between the provider and the client
- ▶ Cloud providers are providing you with the necessary tools to handle these questions



Privacy Protection

- ▶ Data should be protected from unauthorized access
 - ▶ Data encryption
 - ▶ Control who sees and can access what




Privacy Protection

- ▶ Data should be protected from unauthorized access
 - ▶ Data encryption
 - ▶ Control who sees and can access what
- ▶ Identify and classify your data



Privacy Protection

- ▶ Data should be protected from unauthorized access
 - ▶ Data encryption
 - ▶ Control who sees and can access what
- ▶ Identify and classify your data
- ▶ Amazon Macie ! → inventory of Amazon S3 buckets



- ▶ Protect data from unauthorized modification or deletion
- ▶ Challenging in the cloud and especially in a multicloud environment
- ▶ Authorization becomes crucial
 - ▶ Double authentication
 - ▶ Logging



- ▶ Protect data from unauthorized modification or deletion
- ▶ Challenging in the cloud and especially in a multicloud environment
- ▶ Authorization becomes crucial
 - ▶ Double authentication
 - ▶ Logging
- ▶ AWS Key Management Service and EncryptionContext
 - ▶ Additional authenticated data (AAD)
 - ▶ Audit trail
 - ▶ Authorization context




Data Availability: Service Level Agreement (SLA)

- ▶ A Service Level Agreement is a document that defines the quality of service required between a provider and a client
- ▶ **Definition** : an explicit statement of expectations and obligations that exist in a business relationship between two organizations: the service provider and the customer
- ▶ There are several languages of SLA
 - ▶ WS-Agreement by OGF (Open Grid Forum)
 - ▶ WSLA by IBM
- ▶ An SLA is a document of a legal nature
 - ▶ if the SLA is not satisfied, there will be penalties to pay



SLA in the Cloud

- ▶ Often several SLAs because there are several actors involved
- ▶ Ideally, the SLA between a cloud provider and users must be negotiated.
- ▶ Unfortunately, often suppliers offer a single contract for everyone



Service Level Objective (SLO)

- ▶ A SLA is composed of a set of SLOs
- ▶ A SLO specifies the limits of the acceptable
 - ▶ availability > 95%
 - ▶ response time $t < 5$ millisecond
- ▶ A SLA Violation is Made When a SLO Is Not Respected
 - ▶ The penalty is often proportional to the violation
- ▶ Often SLOs are about non-functional requirements



Cloud Non-Functional Requirements

- ▶ **Availability:** an essential requirement, the service must be always present even in case of breakdown -
 - ▶ e.g. 99.99% during work days, 99.9% for nights/weekends)
- ▶ **Reliability: disaster Recovery expectations**
 - ▶ System operation must be consistent
 - ▶ Data must be stored in redundancy (e.g. RAID systems)
 - ▶ Multiple VMs must provide the same tasks
- ▶ **Performance**
 - ▶ Maximum response times
- ▶ **Location of the data:** consistent with local legislation
- ▶ **Portability of the data** (e.g. ability to move data to a different provider)



How to manage infrastructure state effectively ?

Visibility

- ▶ Do you know the state of your infrastructure ? (what, how, who)
- ▶ Can you detect misconfigurations or incident with your service?



How to manage infrastructure state effectively ?

Visibility

- ▶ Do you know the state of your infrastructure ? (what, how, who)
- ▶ Can you detect misconfigurations or incident with your service?

Automation

- ▶ Can you reconfigure your infrastructure to reproduce past configurations or scale up existing ones without a lot of extra manual work?
- ▶ Can you respond to incidents easily or automatically?



How to manage infrastructure state effectively ?

Visibility

- ▶ Do you know the state of your infrastructure ? (what, how, who)
- ▶ Can you detect misconfigurations or incident with your service?

Automation

- ▶ Can you reconfigure your infrastructure to reproduce past configurations or scale up existing ones without a lot of extra manual work?
- ▶ Can you respond to incidents easily or automatically?

Flexibility

- ▶ Can you improve your configurations and scale up easily ?
- ▶ Can you add more complexity using the same tools?
- ▶ Do you share, review, and improve your configurations within your team?



AWS Configuration Management

AWS Console

- ▶ Web interface to control most of the functionalities
- ▶ When should we use the AWS console ?
 - ▶ Read-only usage
 - ▶ Workable for small systems and teams
 - ▶ Rare operations
- ▶ **!** AWS console is not compatible with automation, reproducibility and team communication (documentation and standardization)



AWS Configuration Management

AWS Console

- ▶ Web interface to control most of the functionalities
- ▶ When should we use the AWS console ?
 - ▶ Read-only usage
 - ▶ Workable for small systems and teams
 - ▶ Rare operations
- ▶ **!** AWS console is not compatible with automation, reproducibility and team communication (documentation and standardization)

AWS Command-Line tools (CLI)

- ▶ Most basic way to save and automate AWS operations
- ▶ ++ well-maintained and covers a large proportion of all AWS services
- ▶ Write simple Bash scripts that invoke aws with specific arguments
- ▶ Interactive (not scripted) use : aws-shell



AWS Configuration Management

APIs and SDKs

- ▶ **SDKs** for using AWS APIs are available in most major languages, with Go, iOS, Java, JavaScript, Python, Ruby, and PHP ^a
- ▶ **! Don't use APIs directly** -> use the SDKs for your preferred language to access APIs.

^a(see <https://github.com/donnemartin/awesome-aws>)

Amazon SDK for Python : Boto3

- ▶ Allows to automate operations
- ▶ Contains a variety of APIs that operate at either a high level or a low level
- ▶ When should I use boto ? -> If you find yourself writing bash script with more than 2 CLI.



Visibility

General Visibility

- ▶ Tagging Essential practice to better understand the resource usage
 - ▶ Through automation or convention, you can add tags
 - ▶ For the developer that “owns” that resource
 - ▶ For the product that resource supports
 - ▶ To label lifecycles, such as temporary resources
 - ▶ To distinguish resources with special security or compliance needs
 - ▶ To distinguish production-critical infrastructure
- ! Tags are case sensitive



Visibility

General Visibility

- ▶ Tagging Essential practice to better understand the resource usage
 - ▶ Through automation or convention, you can add tags
 - ▶ For the developer that “owns” that resource
 - ▶ For the product that resource supports
 - ▶ To label lifecycles, such as temporary resources
 - ▶ To distinguish resources with special security or compliance needs
 - ▶ To distinguish production-critical infrastructure
- ! Tags are case sensitive

AWS Services for visibility

- ▶ Access management : IAM
- ▶ Monitoring resources : **Cloud Watch**
- ▶ Monitoring team : Cloud trail



CloudWatch

Monitors resources and applications, captures **logs**, and **sends events**.

- ▶ Standard service for **keeping tabs on AWS resources** :
 - ▶ **EC2** : percentage of allocated EC2 compute units, number CPU credits spent/earned by the instance etc.
 - ▶ **Bucket storage** : amount of data (in bytes), number of objects, number of HTTP requests made to bucket etc.



Monitors resources and applications, captures logs, and sends events.

- ▶ Standard service for **keeping tabs on AWS resources** :
 - ▶ **EC2** : percentage of allocated EC2 compute units, number CPU credits spent/earned by the instance etc.
 - ▶ **Bucket storage** : amount of data (in bytes), number of objects, number of HTTP requests made to bucket etc.
- ▶ Allows to create time based graphs, alarms, and dashboards
 - ▶ Alarms to **trigger notifications** from any given metric.
 - ▶ Trigger SNS notifications, Auto Scaling actions, or EC2 actions.
 - ▶ Publish and share graphs of metrics by creating customizable dashboard views.



CloudWatch

Monitors resources and **applications**, captures **logs**, and **sends events**.

- ▶ Standard service for **keeping tabs on AWS resources** :
 - ▶ **EC2** : percentage of allocated EC2 compute units, number CPU credits spent/earned by the instance etc.
 - ▶ **Bucket storage** : amount of data (in bytes), number of objects, number of HTTP requests made to bucket etc.
- ▶ Allows to create time based graphs, alarms, and dashboards
 - ▶ Alarms to **trigger notifications** from any given metric.
 - ▶ Trigger SNS notifications, Auto Scaling actions, or EC2 actions.
 - ▶ Publish and share graphs of metrics by creating customizable dashboard views.

Limitations

- ▶ Limited to one metric for an alarm.
- ▶ Alarm's notification do not have any contextual detail;
- ▶ Data about metrics is kept in CloudWatch for 15 months.



CloudWatch

CloudWatch Events

- ▶ Mechanism to **automate actions** in various services on AWS.
- ▶ Can create **event rules** from instance states, AWS APIs, Auto Scaling, Run commands, deployments or time-based schedules.
- ▶ **Triggered events** can invoke Lambda functions, send messages, or perform instance actions.



CloudWatch

CloudWatch Events

- ▶ Mechanism to **automate actions** in various services on AWS.
- ▶ Can create **event rules** from instance states, AWS APIs, Auto Scaling, Run commands, deployments or time-based schedules.
- ▶ **Triggered events** can invoke Lambda functions, send messages, or perform instance actions.

CloudWatch Logs

- ▶ Streaming log storage system
- ▶ A log agent installed on your servers will process logs over time and send them to CloudWatch Logs.
- ▶ Can export logged data to S3 or stream results to other AWS services.
- ▶ CloudWatch Logs can be encrypted using keys managed through KMS.



AWS Billing and Cost Management

- ▶ AWS offers a **free tier of service**, that allows very limited usage of resources at no cost.
- ▶ **Pay your AWS bill, monitor** your usage and **budget** your costs. Provides features that you can use to
 - ▶ Estimate and plan your AWS costs
 - ▶ Receive alerts if your costs exceed a threshold that you set
 - ▶ Assess your biggest investments in AWS resources
- ▶ Set **billing alerts** to be notified of unexpected costs

Cost Explorer

- ▶ Tool that enables you to view and analyze your costs and usage.
- ▶ View data for up to the last 13 months
- ▶ Forecast for the next 3 months
- For large accounts, the AWS billing console can **time out** or be **too slow** to use.



Other tools for cost management

Tools

- ▶ Teevity Ice (Netflix) - Docker Ice
- ▶ Security Monkey and Cloud Custodian.
- ▶ AWS Simple Monthly Calculator to get an estimate of usage charges for AWS based on your actual usage of AWS services.


Third-Party Services

- ▶ Cloudability
- ▶ CloudHealth Technologies
- ▶ ParkMyCloud.
- ▶ **!** Some of these charge a percentage of your bill



The Quality Management in the Cloud

► Three visions of the quality



Quality of Business (QoBiz)

- ▶ A provider-centric quality metric
- ▶ It expresses the profit of supplier
 - ▶ service price
 - ▶ revenue per user
 - ▶ revenue per transaction
 - ▶ budget
- ▶ A good QoBiz expresses the prosperity of the provider
- ▶ Decides the pricing policy of the suppliers: rate per hour, rate per month, etc.



Quality of Experience (QoE)

- ▶ According to subjective tests
- ▶ QoS is not sufficient to ensure user satisfaction
- ▶ Quality of experience (QoE) is a subjective metric that expresses quality of service as perceived by the user
- ▶ QoE is influenced by several factors:
 - ▶ The context of use
 - ▶ The personality and preferences of the user
 - ▶ The system used to consume the service

